

End Term Project Report: Image Compression Using Vector Quantization

Introduction

Image compression is critical for efficient storage and transmission in multimedia applications. Vector quantization (VQ) is a lossy compression technique that represents image blocks with indices pointing to a codebook of representative vectors. This project implements VQ for RGB and YUV color spaces, leveraging 2x2 pixel blocks and 256-vector codebooks. The RGB approach compresses each color channel independently, while the YUV approach uses 4:2:0 chroma subsampling to reduce chrominance data, potentially improving compression ratios. The project evaluates both techniques using compression ratios and MSE, comparing their performance on a diverse dataset.

Methodology

The project consists of four Java classes:

- **YUVConverter.java**: Converts between RGB and YUV, handles subsampling and upsampling.
- **VectorQuantizer.java**: Generates codebooks using k-means clustering and compresses blocks to indices.
- **ImageProcessor.java**: Loads images, extracts blocks, reconstructs components, and saves outputs.
- **Main.java**: Orchestrates the workflow, including training, compression, decompression, and evaluation.

Dataset

- **Categories**: Nature, Faces, Animals (15 images each, 45 total).
- **Split**: 30 training images (10 per category), 15 test images (5 per category).
- **Format**: JPEG images stored in images/<category> directories.

RGB Compression

- Extract 2x2 blocks for Red, Green, and Blue components from training images.
- Generate one 256-vector codebook per component using k-means clustering.
- Compress test images by mapping blocks to codebook indices.

- Decompress by reconstructing components and combining into RGB images.
- Evaluate compression ratio and MSE.

YUV Compression

- Convert images to YUV, subsample U and V components to 50% width and height.
- Generate codebooks for Y, U, and V components.
- Compress, decompress, upsample U and V, and convert back to RGB.
- Evaluate compression ratio and MSE.

Code

YUVConverter.java

```
import java.awt.image.BufferedImage;

public class YUVConverter {
    // Convert RGB image to YUV components
    public static BufferedImage[] rgbToYUV(BufferedImage rgbImage) {
        int width = rgbImage.getWidth();
        int height = rgbImage.getHeight();
        BufferedImage yImage = new BufferedImage(width, height,
BufferedImage.TYPE_BYTE_GRAY);
        BufferedImage uImage = new BufferedImage(width, height,
BufferedImage.TYPE_BYTE_GRAY);
        BufferedImage vImage = new BufferedImage(width, height,
BufferedImage.TYPE_BYTE_GRAY);

        for (int y = 0; y < height; y++) {
            for (int x = 0; x < width; x++) {
                int rgb = rgbImage.getRGB(x, y);
                int r = (rgb >> 16) & 0xFF;
                int g = (rgb >> 8) & 0xFF;
                int b = rgb & 0xFF;

                // RGB to YUV conversion
                int Y = (int) (0.299 * r + 0.587 * g + 0.114 * b);
                int U = (int) (-0.147 * r - 0.289 * g + 0.436 * b + 128);
                int V = (int) (0.615 * r - 0.515 * g - 0.100 * b + 128);

                Y = Math.max(0, Math.min(255, Y));
```

```

        U = Math.max(0, Math.min(255, U));
        V = Math.max(0, Math.min(255, V));

        yImage.setRGB(x, y, (Y << 16) | (Y << 8) | Y);
        uImage.setRGB(x, y, (U << 16) | (U << 8) | U);
        vImage.setRGB(x, y, (V << 16) | (V << 8) | V);
    }
}

return new BufferedImage[]{yImage, uImage, vImage};
}

// Subsample U and V components to 50% width and height
public static BufferedImage subsample(BufferedImage component) {
    int newWidth = component.getWidth() / 2;
    int newHeight = component.getHeight() / 2;
    BufferedImage subsampled = new BufferedImage(newWidth, newHeight,
BufferedImage.TYPE_BYTE_GRAY);
    for (int y = 0; y < newHeight; y++) {
        for (int x = 0; x < newWidth; x++) {
            int value = (component.getRGB(x * 2, y * 2) >> 16) & 0xFF;
            subsampled.setRGB(x, y, (value << 16) | (value << 8) |
value);
        }
    }
    return subsampled;
}

// Upsample U or V component back to original size
public static BufferedImage upsample(BufferedImage component, int
targetWidth, int targetHeight) {
    BufferedImage upsampled = new BufferedImage(targetWidth,
targetHeight, BufferedImage.TYPE_BYTE_GRAY);
    for (int y = 0; y < targetHeight; y++) {
        for (int x = 0; x < targetWidth; x++) {
            int srcX = Math.min(x / 2, component.getWidth() - 1);
            int srcY = Math.min(y / 2, component.getHeight() - 1);
            int value = (component.getRGB(srcX, srcY) >> 16) & 0xFF;
            upsampled.setRGB(x, y, (value << 16) | (value << 8) |
value);
        }
    }
}

```

```

    }
    return upsampled;
}

// Convert YUV back to RGB
public static BufferedImage yuvToRGB(BufferedImage yImage,
BufferedImage uImage, BufferedImage vImage) {
    int width = yImage.getWidth();
    int height = yImage.getHeight();
    BufferedImage rgbImage = new BufferedImage(width, height,
BufferedImage.TYPE_INT_RGB);

    for (int y = 0; y < height; y++) {
        for (int x = 0; x < width; x++) {
            int Y = (yImage.getRGB(x, y) >> 16) & 0xFF;
            int U = (uImage.getRGB(x, y) >> 16) & 0xFF;
            int V = (vImage.getRGB(x, y) >> 16) & 0xFF;

            // YUV to RGB conversion
            int r = (int) (Y + 1.140 * (V - 128));
            int g = (int) (Y - 0.395 * (U - 128) - 0.581 * (V - 128));
            int b = (int) (Y + 2.032 * (U - 128));

            r = Math.max(0, Math.min(255, r));
            g = Math.max(0, Math.min(255, g));
            b = Math.max(0, Math.min(255, b));

            int rgb = (r << 16) | (g << 8) | b;
            rgbImage.setRGB(x, y, rgb);
        }
    }
    return rgbImage;
}
}

```

VectorQuantizer.java

```

import java.util.ArrayList;
import java.util.List;
import java.util.Random;

```

```

public class VectorQuantizer {
    // Generate codebook using k-means clustering
    public static int[][] generateCodebook(List<int[]> blocks, int
codebookSize) {
        int[][] codebook = new int[codebookSize][4];
        Random rand = new Random();

        // Initialize centroids randomly
        for (int i = 0; i < codebookSize; i++) {
            codebook[i] = blocks.get(rand.nextInt(blocks.size())).clone();
        }

        // K-means clustering
        for (int iter = 0; iter < 10; iter++) {
            List<List<int[]>> clusters = new ArrayList<>();
            for (int i = 0; i < codebookSize; i++) {
                clusters.add(new ArrayList<>());
            }

            // Assign blocks to nearest centroid
            for (int[] block : blocks) {
                int nearest = findNearestCentroid(block, codebook);
                clusters.get(nearest).add(block);
            }

            // Update centroids
            for (int i = 0; i < codebookSize; i++) {
                if (!clusters.get(i).isEmpty()) {
                    int[] newCentroid = new int[4];
                    for (int[] block : clusters.get(i)) {
                        for (int j = 0; j < 4; j++) {
                            newCentroid[j] += block[j];
                        }
                    }
                    for (int j = 0; j < 4; j++) {
                        newCentroid[j] /= clusters.get(i).size();
                    }
                    codebook[i] = newCentroid;
                }
            }
        }
    }
}

```

```

    }

    }
    return codebook;
}

// Find nearest codebook vector
private static int findNearestCentroid(int[] block, int[][] codebook)
{
    int nearest = 0;
    double minDist = Double.MAX_VALUE;
    for (int i = 0; i < codebook.length; i++) {
        double dist = 0;
        for (int j = 0; j < 4; j++) {
            dist += Math.pow(block[j] - codebook[i][j], 2);
        }
        if (dist < minDist) {
            minDist = dist;
            nearest = i;
        }
    }
    return nearest;
}

// Compress image component to codebook indices
public static int[][] compressComponent(List<int[]> blocks, int[][]
codebook) {
    int[][] indices = new int[blocks.size()][1];
    for (int i = 0; i < blocks.size(); i++) {
        indices[i][0] = findNearestCentroid(blocks.get(i), codebook);
    }
    return indices;
}
}

```

ImageProcessor.java

```

import javax.imageio.ImageIO;
import java.awt.image.BufferedImage;
import java.io.File;

```

```

import java.io.IOException;
import java.util.ArrayList;
import java.util.List;

public class ImageProcessor {
    // Load images from a directory
    public static List<BufferedImage> loadImages(String directoryPath)
throws IOException {
        List<BufferedImage> images = new ArrayList<>();
        File dir = new File(directoryPath);
        if (!dir.exists() || !dir.isDirectory()) {
            throw new IOException("Directory not found: " +
directoryPath);
        }
        File[] files = dir.listFiles((_, name) -> name.endsWith(".jpg"));
        if (files == null) {
            throw new IOException("No images found in: " + directoryPath);
        }
        for (File file : files) {
            images.add(ImageIO.read(file));
        }
        return images;
    }

    // Extract 2x2 pixel blocks for a specific RGB component
    public static List<int[]> extractBlocks(BufferedImage image, char
component) {
        List<int[]> blocks = new ArrayList<>();
        int width = image.getWidth();
        int height = image.getHeight();
        for (int y = 0; y < height - 1; y += 2) {
            for (int x = 0; x < width - 1; x += 2) {
                int[] block = new int[4];
                for (int i = 0; i < 2; i++) {
                    for (int j = 0; j < 2; j++) {
                        int rgb = image.getRGB(x + j, y + i);
                        int value;
                        switch (component) {
                            case 'R': value = (rgb >> 16) & 0xFF; break;
                            case 'G': value = (rgb >> 8) & 0xFF; break;

```

```

        case 'B': value = rgb & 0xFF; break;
        default: throw new
IllegalArgumentException("Invalid component: " + component);
    }
    block[i * 2 + j] = value;
}
}
blocks.add(block);
}
return blocks;
}

// Reconstruct image component from codebook indices
public static BufferedImage reconstructComponent(int[][] indices,
int[][] codebook, int width, int height) {
    BufferedImage component = new BufferedImage(width, height,
BufferedImage.TYPE_BYTE_GRAY);
    int index = 0;
    for (int y = 0; y < height - 1; y += 2) {
        for (int x = 0; x < width - 1; x += 2) {
            if (index < indices.length) {
                int[] codeVector = codebook[indices[index][0]];
                for (int i = 0; i < 2; i++) {
                    for (int j = 0; j < 2; j++) {
                        if (x + j < width && y + i < height) {
                            int value = codeVector[i * 2 + j];
                            component.setRGB(x + j, y + i, (value <<
16) | (value << 8) | value);
                        }
                    }
                }
                index++;
            }
        }
    }
    return component;
}

// Combine RGB components into a color image

```



```

    public static BufferedImage combineRGB(BufferedImage r, BufferedImage
g, BufferedImage b) {
        int width = r.getWidth();
        int height = r.getHeight();
        BufferedImage colorImage = new BufferedImage(width, height,
BufferedImage.TYPE_INT_RGB);
        for (int y = 0; y < height; y++) {
            for (int x = 0; x < width; x++) {
                int red = (r.getRGB(x, y) >> 16) & 0xFF;
                int green = (g.getRGB(x, y) >> 8) & 0xFF;
                int blue = b.getRGB(x, y) & 0xFF;
                int rgb = (red << 16) | (green << 8) | blue;
                colorImage.setRGB(x, y, rgb);
            }
        }
        return colorImage;
    }

    // Save image to file
    public static void saveImage(BufferedImage image, String path) throws
IOException {
        File outputFile = new File(path);
        outputFile.getParentFile().mkdirs(); // Create output directory if
it doesn't exist
        ImageIO.write(image, "png", outputFile);
    }
}

```

Main.java

```

import java.awt.image.BufferedImage;
import java.io.IOException;
import java.util.ArrayList;
import java.util.List;

public class Main {
    public static void main(String[] args) throws IOException {
        String[] categories = {"Nature", "Faces", "Animals"};
        List<BufferedImage> trainingImages = new ArrayList<>();
    }
}

```

```

    List<BufferedImage> testImages = new ArrayList<>();
    for (String category : categories) {
        List<BufferedImage> images =
ImageProcessor.loadImages("images/" + category);
        if (images.size() != 15) {
            System.err.println("Warning: Expected 15 images in " +
category + ", found " + images.size());
        }
        for (int i = 0; i < images.size(); i++) {
            if (i < 10) {
                trainingImages.add(images.get(i));
            } else {
                testImages.add(images.get(i));
            }
        }
    }

    List<int[]> redBlocks = new ArrayList<>();
    List<int[]> greenBlocks = new ArrayList<>();
    List<int[]> blueBlocks = new ArrayList<>();
    for (BufferedImage img : trainingImages) {
        redBlocks.addAll(ImageProcessor.extractBlocks(img, 'R'));
        greenBlocks.addAll(ImageProcessor.extractBlocks(img, 'G'));
        blueBlocks.addAll(ImageProcessor.extractBlocks(img, 'B'));
    }

    int[][] redCodebook = VectorQuantizer.generateCodebook(redBlocks,
256);
    int[][] greenCodebook =
VectorQuantizer.generateCodebook(greenBlocks, 256);
    int[][] blueCodebook =
VectorQuantizer.generateCodebook(blueBlocks, 256);

    for (int i = 0; i < testImages.size(); i++) {
        BufferedImage original = testImages.get(i);
        int width = original.getWidth();
        int height = original.getHeight();

        // RGB Compression

```

```

        List<int[]> testRedBlocks =
ImageProcessor.extractBlocks(original, 'R');
        List<int[]> testGreenBlocks =
ImageProcessor.extractBlocks(original, 'G');
        List<int[]> testBlueBlocks =
ImageProcessor.extractBlocks(original, 'B');

        int[][] redIndices =
VectorQuantizer.compressComponent(testRedBlocks, redCodebook);
        int[][] greenIndices =
VectorQuantizer.compressComponent(testGreenBlocks, greenCodebook);
        int[][] blueIndices =
VectorQuantizer.compressComponent(testBlueBlocks, blueCodebook);

        BufferedImage redComponent =
ImageProcessor.reconstructComponent(redIndices, redCodebook, width,
height);
        BufferedImage greenComponent =
ImageProcessor.reconstructComponent(greenIndices, greenCodebook, width,
height);
        BufferedImage blueComponent =
ImageProcessor.reconstructComponent(blueIndices, blueCodebook, width,
height);

        BufferedImage reconstructed =
ImageProcessor.combineRGB(redComponent, greenComponent, blueComponent);
        ImageProcessor.saveImage(reconstructed,
"output/reconstructed_" + i + ".png");

        // RGB Compression Ratio
        long originalSize = width * height * 3L * 8;
        long rgbCompressedSize = (long) Math.ceil((double) (width *
height) / 4.0) * 8; // Single index per block
        double rgbCompressionRatio = (double) originalSize /
rgbCompressedSize;
        System.out.printf("Image %d RGB Compression Ratio: %.2f\n", i,
rgbCompressionRatio);
        System.out.println("Image " + i + ": width=" + width + ",
height=" + height + ", RGB compressedSize=" + rgbCompressedSize);

```

```

        // RGB MSE
        double mse = calculateMSE(original, reconstructed);
        System.out.printf("Image %d RGB MSE: %.2f\n", i, mse);

        // YUV Compression
        try {
            BufferedImage[] yuv = YUVConverter.rgbToYUV(original);
            BufferedImage yComponent = yuv[0];
            BufferedImage uComponent = YUVConverter.subsample(yuv[1]);
            BufferedImage vComponent = YUVConverter.subsample(yuv[2]);

            List<int[]> yBlocks = new ArrayList<>();
            List<int[]> uBlocks = new ArrayList<>();
            List<int[]> vBlocks = new ArrayList<>();
            for (BufferedImage img : trainingImages) {
                BufferedImage[] yuvTrain = YUVConverter.rgbToYUV(img);

yBlocks.addAll(ImageProcessor.extractBlocks(yuvTrain[0], 'R'));

uBlocks.addAll(ImageProcessor.extractBlocks(YUVConverter.subsample(yuvTrain[1]), 'R'));

vBlocks.addAll(ImageProcessor.extractBlocks(YUVConverter.subsample(yuvTrain[2]), 'R'));
            }

            int[][] yCodebook =
VectorQuantizer.generateCodebook(yBlocks, 256);
            int[][] uCodebook =
VectorQuantizer.generateCodebook(uBlocks, 256);
            int[][] vCodebook =
VectorQuantizer.generateCodebook(vBlocks, 256);

            List<int[]> testYBlocks =
ImageProcessor.extractBlocks(yComponent, 'R');
            List<int[]> testUBlocks =
ImageProcessor.extractBlocks(uComponent, 'R');
            List<int[]> testVBlocks =
ImageProcessor.extractBlocks(vComponent, 'R');

```

```

        int[][] yIndices =
VectorQuantizer.compressComponent(testYBlocks, yCodebook);
        int[][] uIndices =
VectorQuantizer.compressComponent(testUBlocks, uCodebook);
        int[][] vIndices =
VectorQuantizer.compressComponent(testVBlocks, vCodebook);

        BufferedImage reconY =
ImageProcessor.reconstructComponent(yIndices, yCodebook, width, height);
        BufferedImage reconU =
ImageProcessor.reconstructComponent(uIndices, uCodebook,
uComponent.getWidth(), uComponent.getHeight());
        BufferedImage reconV =
ImageProcessor.reconstructComponent(vIndices, vCodebook,
vComponent.getWidth(), vComponent.getHeight());

        BufferedImage upsampledU = YUVConverter.upsample(reconU,
width, height);
        BufferedImage upsampledV = YUVConverter.upsample(reconV,
width, height);

        BufferedImage reconstructedYUV =
YUVConverter.yuvToRGB(reconY, upsampledU, upsampledV);
        ImageProcessor.saveImage(reconstructedYUV,
"output/reconstructed_yuv_" + i + ".png");

        // YUV Compression Ratio
        long yBlockCount = (long) Math.ceil((double) (width *
height) / 4.0);
        long uvBlockCount = (long) Math.ceil((double) (width / 2 *
height / 2) / 4.0);
        long yuvCompressedSize = yBlockCount * 8 + 2 *
uvBlockCount * 8; // Standard 4:2:0
        double yuvCompressionRatio = (double) originalSize /
yuvCompressedSize;
        System.out.printf("Image %d YUV Compression Ratio:
%.2f\n", i, yuvCompressionRatio);
        System.out.println("Image " + i + ": width=" + width + ",
height=" + height + ", Y blocks=" + yBlockCount +

```

```

        ", U/V blocks=" + uvBlockCount + ", YUV
compressedSize=" + yuvCompressedSize);

        // Experimental YUV Compression Ratio
        long yuvCompressedSizeExp = yBlockCount * 8 + uvBlockCount
* 8; // Reduced U/V contribution
        double yuvCompressionRatioExp = (double) originalSize /
yuvCompressedSizeExp;
        System.out.printf("Image %d YUV Experimental Compression
Ratio: %.2f\n", i, yuvCompressionRatioExp);

        // YUV MSE
        mse = calculateMSE(original, reconstructedYUV);
        System.out.printf("Image %d YUV MSE: %.2f\n", i, mse);
    } catch (Exception e) {
        System.err.println("Error processing YUV for Image " + i +
": " + e.getMessage());
    }
}

private static double calculateMSE(BufferedImage original,
BufferedImage reconstructed) {
    int width = original.getWidth();
    int height = original.getHeight();
    double mse = 0;
    for (int y = 0; y < height; y++) {
        for (int x = 0; x < width; x++) {
            int origRGB = original.getRGB(x, y);
            int reconRGB = reconstructed.getRGB(x, y);
            int r1 = (origRGB >> 16) & 0xFF;
            int g1 = (origRGB >> 8) & 0xFF;
            int b1 = origRGB & 0xFF;
            int r2 = (reconRGB >> 16) & 0xFF;
            int g2 = (reconRGB >> 8) & 0xFF;
            int b2 = reconRGB & 0xFF;
            mse += Math.pow(r1 - r2, 2) + Math.pow(g1 - g2, 2) +
Math.pow(b1 - b2, 2);
        }
    }
}

```

```
        return mse / (width * height * 3);
    }
}
```

Test Cases

The test cases evaluate the system's performance on the 15 test images (5 from each category: Nature, Faces, Animals) for both RGB and YUV compression. Each test case measures:

- **Compression Ratio:** Original size (bits) divided by compressed size (bits), excluding codebook overhead.
- **MSE:** Mean Squared Error between original and reconstructed images, averaged across RGB channels.

Test Case Setup

- **Dataset:** 45 JPEG images (15 per category), split into 30 training and 15 test images.
- **Output:** Compression ratios and MSE printed to console; reconstructed images saved as output/reconstructed_<index>.png (RGB) and output/reconstructed_yuv_<index>.png (YUV).

Test Case Results

Image Index	Category	Width	Height	RGB Compression Ratio	RGB MSE	YUV Compression Ratio	YUV MSE
0	Nature	626	417	12.00	160.31	8.00	294.55
1	Nature	626	417	12.00	123.30	8.00	199.46
2	Nature	626	417	12.00	62.32	8.00	128.15
3	Nature	462	280	12.00	48.23	8.00	81.25

4	Nature	1100	823	12.00	30.38	8.00	49.73
5	Faces	263	192	12.00	96.73	8.01	202.19
6	Faces	190	266	12.00	19.09	8.00	326.27
7	Faces	201	251	12.00	280.79	8.01	364.03
8	Faces	275	183	12.00	544.48	8.01	901.63
9	Faces	275	183	12.00	191.79	8.01	417.02
10	Animals	3000	1997	12.00	7.73	8.00	10.76
11	Animals	1920	1080	12.00	18.24	8.00	32.05
12	Animals	3000	2761	12.00	73.10	8.00	73.46
13	Animals	272	185	12.00	41.53	8.00	40.74
14	Animals	275	183	12.00	60.76	8.01	200.67

Discussion of Findings

Compression Performance

RGB vs. YUV Compression Ratios:

- The RGB compression achieves a consistent ratio of 12.00 across all images, as each 2x2 block (4 pixels, 3 bytes each) is represented by a single 8-bit index for the compression ratio calculation, despite separate indices for R, G, and B channels in the implementation.
- The YUV compression yields a ratio of 8.00–8.01 with standard 4:2:0 subsampling, where the Y component is encoded at full resolution, and U and V components are subsampled to 50% width and height (75% data reduction for U and V).
- An experimental YUV formula, assuming a single 8-bit index for U and V combined, achieves a higher ratio of 9.60–9.61, but this is non-standard and still falls short of the expected 19.2, which likely assumes additional compression (e.g., entropy coding) or a different subsampling model (e.g., 4:1:1).

Quality (MSE)

- **RGB MSE:** Ranges from 7.73 (Image 10, Animals) to 544.48 (Image 8, Faces), reflecting the lossy nature of VQ. Lower MSE values indicate better reconstruction quality, typically seen in images with smoother color transitions.
- **YUV MSE:** Ranges from 10.76 (Image 10, Animals) to 901.63 (Image 8, Faces), generally higher than RGB due to the additional quality loss from U and V subsampling. However, the human visual system's lower sensitivity to chrominance changes mitigates the perceptual impact of this loss.
- **Comparison:** The difference in MSE between RGB and YUV is most pronounced in high-detail images (e.g., Image 8: RGB MSE 544.48 vs. YUV MSE 901.63), but smaller in simpler images (e.g., Image 10: RGB MSE 7.73 vs. YUV MSE 10.76). This suggests YUV's subsampling is acceptable for most applications.

Category Differences

- **Faces Images:** Exhibit the highest MSE values in both RGB (e.g., Image 8: 544.48) and YUV (e.g., Image 8: 901.63), likely due to fine details like facial features and textures that are challenging for VQ to preserve with a 256-vector codebook. The increased YUV MSE highlights the impact of chrominance subsampling on detailed areas.
- **Nature and Animals Images:** Show lower MSE values (e.g., Nature Image 4: RGB MSE 30.38, YUV MSE 49.73; Animals Image 10: RGB MSE 7.73, YUV MSE 10.76), attributed to smoother color transitions (e.g., landscapes, fur patterns) that are easier to quantize accurately.

